

Prajwal Kulkarni

Dallas, TX — (763) 306-5686 — prajwalkp.work@gmail.com — linkedin.com/in/prajwal-kp — kp-prajwal.github.io

EDUCATION

The University of Texas at Dallas

Master of Science, Business Analytics & Artificial Intelligence

GPA: 3.8/4.0

TECHNICAL SKILLS

Certifications: AWS Databricks Platform Architect, Snowflake Associate: Platform, Azure Data Fundamentals

Programming & Query: Python, SQL (BigQuery, NoSQL, DuckDB, DynamoDB), R, Bash/Shell, Git

Data & Orchestration: PySpark, Apache Airflow, Apache Kafka, Hadoop, Hive, GitHub Actions

Cloud & Platforms: Databricks, Azure (ADLS, ADF, Synapse, Event Hub), AWS (EC2, S3, Lambda, SageMaker, EMR, Glue), Snowflake

Analytics & Governance: Power BI, Tableau, Docker, Unity Catalog, Grafana, QuickSight, Jira, Stonebranch

EXPERIENCE

Albertsons Companies — AI Data Engineer

Dallas, TX | Jun 2025 – Aug 2026

- Architected a greenfield finance data platform using medallion architecture on Databricks across GCP and Azure, producing 40 standardized P&L metrics for financial forecasting across 2,500+ U.S. stores.
- Productionized 40+ PySpark/SQL Databricks pipelines processing 1B+ records per P&L line, with full support for historical backfills and weekly/quarterly incremental workloads, eliminating manual reconciliation that previously consumed finance analyst hours each cycle.
- Reduced compute costs on billion-row pipelines by redesigning data models with partitioning and materialized views, eliminating redundant full-table scans across critical finance workloads.
- Built executive Power BI dashboards for AOP variance, allocation, and corporate forecast analysis, adopted as the weekly decision-making tool by merchandising and finance leadership across store clusters.
- Shipped two GenAI capabilities to production: an anomaly detection layer that surfaces financial KPI outliers before weekly reviews, and a natural-language variance summarizer that auto-generates root-cause narratives.

GenerativeProduct.io — GenAI Data Engineer

Dallas, TX | May 2023 – Apr 2025

- Built an LLM-powered interview simulation platform from zero to 300+ active users using LangChain and open-source transformer models, handling conversational flows without reliance on managed LLM APIs.
- Designed a RAG system using embedding models and vector search to dynamically select and rank interview questions by candidate profile, increasing average session duration by 25% versus static question delivery.
- Led GenAI experimentation by designing A/B tests on prompt engineering, retrieval strategies, and response ranking; evaluated model performance using statistical testing across 5,000+ sessions to guide product decisions.

ABB Ltd. — Data Analytics Engineer

Bangalore, India | Jul 2021 – Jul 2023

- Designed Azure-based batch pipelines using Spark and Data Factory to ingest 50M+ IoT sensor records per day from distributed industrial equipment, creating the data foundation for a predictive maintenance program.
- Operationalized end-to-end ML pipelines covering feature engineering, training, scheduled retraining, and real-time inference for equipment health scoring.
- Trained Random Forest time-series forecasting models that cut anomaly detection false positives by 75%, validated using ROC-AUC, MAE, and RMSE on held-out production windows.
- Implemented CI/CD for ML workflows with GitHub Actions and Azure DevOps, enabling automated testing, version-controlled model artifacts, and zero-downtime rollouts across dev, staging, and production environments.

PROJECTS

Conversational AI Agent for UT Dallas – LLM & NLP

- Deployed a conversational AI chatbot for 5,000+ users using transformer-based LLMs with voice input.
- Implemented retrieval-based intent handling and confidence scoring to improve response accuracy.